

## 模型训练

本notebook 是在沃特·老师您在苏大“24Fall”专业实践小组作业的模型训练部分。  
本部分实现并训练了三个模型用于CIFar10 数据集的分类：

- alexnet
- mobilenetv3
- shufflenetv2

## 导入库

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import torch
import torchvision
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
import random
from torch.utils.data import DataLoader
import string
import os

In [2]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
device

Out[2]: device(type='cuda')
```

## 模型定义

### 1. AlexNet

```
In [3]: class AlexNet(nn.Module):
    def __init__(self, num_classes=1000, dropout=0.5):
        super(AlexNet, self).__init__()
        self.features = nn.Sequential(
            # 论文中的输出通道数为96, pytorch官方为64
            # nn.Conv2d(in_channels=3, out_channels=64, kernel_size=11, stride=4, padding=2),
            nn.Conv2d(in_channels=3, out_channels=96, kernel_size=11, stride=4, padding=2),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
            # 论文中的输出通道数为256, pytorch官方为192
            # nn.Conv2d(in_channels=64, out_channels=192, kernel_size=5, padding=2),
            nn.Conv2d(in_channels=96, out_channels=256, kernel_size=5, padding=2),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
            # nn.Conv2d(in_channels=192, out_channels=384, kernel_size=3, padding=1),
            nn.Conv2d(in_channels=256, out_channels=384, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(in_channels=384, out_channels=256, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2)
        )
        # 这一操作是为了保证特征提取后的特征图大小为 6*6, 使得网络可以接受224x224以外尺寸的图片
        self.avgpool = nn.AdaptiveAvgPool2d(output_size=(6, 6))
        self.classifier = nn.Sequential(
            nn.Dropout(p=dropout),
            nn.Linear(256 * 6 * 6 = 5, 4096),
            nn.ReLU(inplace=True),
            nn.Dropout(p=dropout),
            nn.Linear(4096, 4096),
            nn.ReLU(inplace=True),
            nn.Linear(4096, num_classes)
        )

    def forward(self, x):
        # 提取图像特征
        x = self.features(x)
        x = self.avgpool(x)
        x = torch.flatten(x, start_dim=1)
        # 进行图像分类
        x = self.classifier(x)
        return x

In [4]: model_AlexNet = AlexNet(num_classes=10)
inputs = torch.randn(1, 3, 224, 224)
out = model_AlexNet(inputs)
print(out.shape)

torch.Size([1, 10])
```

### 2. Mobilenetv3

```
In [5]: class HardSwish(nn.Module):
    def __init__(self, inplace=True):
        super(HardSwish, self).__init__()
        self.relu6 = nn.ReLU6(inplace=inplace)

    def forward(self, x):
        return x * self.relu6(x*3)/6

class ConvBNActivation(nn.Module):
    def __init__(self, in_channel, out_channel, kernel_size, stride, groups, activate):
        padding = (kernel_size - 1) // 2
        super(ConvBNActivation, self).__init__()
        nn.Conv2d(in_channels=in_channel, out_channels=out_channel, kernel_size=kernel_size, stride=stride, padding=padding, groups=groups, bias=False)
        nn.ReLU(inplace=True) if activate == 'relu' else HardSwish()
    )

class SqueezeAndExcite(nn.Module):
    def __init__(self, in_channel, out_channel, divide=4):
        super(SqueezeAndExcite, self).__init__()
        mid_channel = in_channel // divide
        self.pool = nn.AdaptiveAvgPool2d([1, 1])
        self.SBconv = nn.Sequential(
            nn.Linear(in_features=in_channel, out_features=mid_channel),
            nn.ReLU(inplace=True),
            HardSwish(),
        )

    def forward(self, x):
        b, c, h, w = x.size()
        out = self.pool(x)
        out = torch.flatten(out, start_dim=1)
        out = self.SBconv(out)
        out = self.relu6(out)
        if self.use_shortcut:
            return x + out
        return out

class SEInvertBottleneck(nn.Module):
    def __init__(self, in_channel, mid_channel, out_channel, kernel_size, use_se, activate, stride):
        super(SEInvertBottleneck, self).__init__()
        self.use_shortcut = stride == 1 and in_channel == out_channel
        self.use_se = use_se

        self.conv = ConvBNActivation(in_channel=in_channel, out_channel=mid_channel, kernel_size=1, stride=1, padding=0, bias=False)
        self.depth_conv = ConvBNActivation(in_channel=mid_channel, out_channel=mid_channel, kernel_size=kernel_size, stride=1, padding=0, bias=False)
        self.point_conv = ConvBNActivation(in_channel=mid_channel, out_channel=out_channel, kernel_size=1, stride=1, padding=0, bias=False)

    def forward(self, x):
        out = self.conv(x)
        out = self.depth_conv(out)
        if self.use_se:
            out = self.SBconv(out)
        out = self.point_conv(out)
        if self.use_shortcut:
            return x + out
        return out

class MobileNetV3(nn.Module):
    def __init__(self, num_classes=1000, type='large'):
        super(MobileNetV3, self).__init__()
        self.type = type

        self.first_conv = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=16, kernel_size=3, stride=2, padding=1, bias=False),
            HardSwish(),
        )

        if self.type == 'large':
            self.large_bottleneck = nn.Sequential(
                SEInvertBottleneck(in_channel=16, mid_channel=16, out_channel=16, kernel_size=3, use_se=True, activate=HardSwish(), stride=1),
                SEInvertBottleneck(in_channel=16, mid_channel=72, out_channel=24, kernel_size=3, use_se=True, activate=HardSwish(), stride=2),
                SEInvertBottleneck(in_channel=24, mid_channel=72, out_channel=24, kernel_size=3, use_se=True, activate=HardSwish(), stride=1),
                SEInvertBottleneck(in_channel=24, mid_channel=120, out_channel=40, kernel_size=3, use_se=True, activate=HardSwish(), stride=2),
                SEInvertBottleneck(in_channel=40, mid_channel=120, out_channel=40, kernel_size=3, use_se=True, activate=HardSwish(), stride=1),
                SEInvertBottleneck(in_channel=40, mid_channel=240, out_channel=80, kernel_size=3, use_se=True, activate=HardSwish(), stride=2),
                SEInvertBottleneck(in_channel=80, mid_channel=240, out_channel=80, kernel_size=3, use_se=True, activate=HardSwish(), stride=1),
                SEInvertBottleneck(in_channel=80, mid_channel=184, out_channel=80, kernel_size=3, use_se=True, activate=HardSwish(), stride=2),
                SEInvertBottleneck(in_channel=80, mid_channel=184, out_channel=80, kernel_size=3, use_se=True, activate=HardSwish(), stride=1),
                SEInvertBottleneck(in_channel=80, mid_channel=480, out_channel=112, kernel_size=3, use_se=True, activate=HardSwish(), stride=2),
                SEInvertBottleneck(in_channel=112, mid_channel=672, out_channel=112, kernel_size=3, use_se=True, activate=HardSwish(), stride=1),
                SEInvertBottleneck(in_channel=112, mid_channel=672, out_channel=160, kernel_size=3, use_se=True, activate=HardSwish(), stride=2),
                SEInvertBottleneck(in_channel=160, mid_channel=160, out_channel=160, kernel_size=3, use_se=True, activate=HardSwish(), stride=1),
            )
            self.large_last_stage = nn.Sequential(
                nn.Conv2d(in_channels=160, out_channels=960, kernel_size=1, stride=1, bias=False),
                nn.BatchNorm2d(960),
                nn.AdaptiveAvgPool2d([1, 1]),
                nn.Conv2d(in_channels=960, out_channels=1280, kernel_size=1, stride=1, bias=False),
                HardSwish(),
            )
        else:
            self.small_bottleneck = nn.Sequential(
                SEInvertBottleneck(in_channel=16, mid_channel=16, out_channel=16, kernel_size=3, use_se=True, activate=HardSwish(), stride=1),
                SEInvertBottleneck(in_channel=16, mid_channel=72, out_channel=24, kernel_size=3, use_se=True, activate=HardSwish(), stride=2),
                SEInvertBottleneck(in_channel=24, mid_channel=72, out_channel=24, kernel_size=3, use_se=True, activate=HardSwish(), stride=1),
                SEInvertBottleneck(in_channel=24, mid_channel=96, out_channel=40, kernel_size=3, use_se=True, activate=HardSwish(), stride=2),
                SEInvertBottleneck(in_channel=40, mid_channel=240, out_channel=40, kernel_size=3, use_se=True, activate=HardSwish(), stride=1),
                SEInvertBottleneck(in_channel=40, mid_channel=120, out_channel=40, kernel_size=3, use_se=True, activate=HardSwish(), stride=2),
                SEInvertBottleneck(in_channel=40, mid_channel=144, out_channel=40, kernel_size=3, use_se=True, activate=HardSwish(), stride=1),
                SEInvertBottleneck(in_channel=40, mid_channel=280, out_channel=96, kernel_size=3, use_se=True, activate=HardSwish(), stride=2),
                SEInvertBottleneck(in_channel=96, mid_channel=576, out_channel=96, kernel_size=3, use_se=True, activate=HardSwish(), stride=1),
                SEInvertBottleneck(in_channel=96, mid_channel=576, out_channel=96, kernel_size=3, use_se=True, activate=HardSwish(), stride=2),
            )
            self.small_last_stage = nn.Sequential(
                nn.Conv2d(in_channels=96, out_channels=576, kernel_size=1, stride=1, bias=False),
                nn.BatchNorm2d(576),
                HardSwish(),
                nn.AdaptiveAvgPool2d([1, 1]),
                nn.Conv2d(in_channels=576, out_channels=1280, kernel_size=1, stride=1, bias=False),
                HardSwish(),
            )

        self.classifier = nn.Sequential(
            nn.Conv2d(in_channels=1280, out_channels=1280, kernel_size=1, stride=1, padding=0, bias=False),
            nn.Linear(in_features=1280, out_features=num_classes),
        )

    # weight init
    for m in self.modules():
        if isinstance(m, nn.Conv2d):
            nn.init.kaiming_normal_(m.weight, mode='fan_out')
            if m.bias is not None:
                nn.init.zeros_(m.bias)
        elif isinstance(m, nn.BatchNorm2d):
            nn.init.ones_(m.weight)
            nn.init.zeros_(m.bias)
        elif isinstance(m, nn.Linear):
            nn.init.normal_(m.weight, mean=0, std=0.01)
            nn.init.zeros_(m.bias)

    def forward(self, x):
        x = self.first_conv(x)
        if self.type == 'large':
            x = self.large_bottleneck(x)
            x = self.large_last_stage(x)
        else:
            x = self.small_bottleneck(x)
            x = self.small_last_stage(x)

        x = torch.flatten(x, start_dim=1)
        x = self.classifier(x)
        return x

In [6]: inputs = torch.randn(1, 3, 224, 224)
model_MobileNetV3 = MobileNetV3(num_classes=10)
out = model_MobileNetV3(inputs)
print(out.shape)

torch.Size([1, 10])
```

### 3. Shufflenetv2

```
In [7]: class ConvBNReLU(nn.Sequential):
    def __init__(self, in_channel, out_channel, kernel_size, stride, groups):
        padding = (kernel_size - 1) // 2
        super(ConvBNReLU, self).__init__()
        nn.Conv2d(in_channels=in_channel, out_channels=out_channel, kernel_size=kernel_size, stride=stride, padding=padding, groups=groups, bias=False)
        nn.ReLU(inplace=True),
        nn.BatchNorm2d(out_channel),
    )

class ConvBN(nn.Sequential):
    def __init__(self, in_channel, out_channel, kernel_size, stride, groups):
        padding = (kernel_size - 1) // 2
        super(ConvBN, self).__init__()
        nn.Conv2d(in_channels=in_channel, out_channels=out_channel, kernel_size=kernel_size, stride=stride, padding=padding, groups=groups, bias=False)
        nn.BatchNorm2d(out_channel),
    )

class HalfSplit(nn.Module):
    """
    将channel split
    """
    def __init__(self, dim=0, first_half=True):
        super(HalfSplit, self).__init__()
        self.first_half = first_half
        self.dim = dim

    def forward(self, x):
        splits = torch.chunk(x, 2, dim=self.dim)
        return splits[0] if self.first_half else splits[1]

class ChannelShuffle(nn.Module):
    def __init__(self, groups):
        super(ChannelShuffle, self).__init__()
        self.groups = groups

    def forward(self, x):
        # Channel shuffle: [N,C,H,W] -> [N,g,C/g,H,W] -> [N,C/g,g,H,W] -> [N,C,H,W]
        batch_size, num_channels, height, width = x.size()
        channels_per_group = num_channels // self.groups
        x = x.view(batch_size, self.groups, channels_per_group, height, width)
        x = torch.transpose(x, dim=0, dim=1).contiguous()
        x = x.view(batch_size, -1, height, width)
        return x

class ShuffleNetUnits(nn.Module):
    def __init__(self, in_channel, out_channel, stride, groups):
        super(ShuffleNetUnits, self).__init__()
        self.stride = stride
        if self.stride > 1:
            mid_channel = out_channel // 2
            in_channel = mid_channel
            self.first_split = HalfSplit(dim=1, first_half=True)
            self.second_split = HalfSplit(dim=1, first_half=False)
        # 论文中f(g,3,d) 中的右半部分
        self.bottleneck = nn.Sequential(
            # 3x3 Conv
            ConvBNReLU(in_channel=in_channel, out_channel=mid_channel, kernel_size=1, stride=1, groups=1),
            # 3x3 DWConv
            ConvBN(in_channel=mid_channel, out_channel=mid_channel, kernel_size=3, stride=stride, groups=1),
            # 3x3 Conv
            ConvBNReLU(in_channel=mid_channel, out_channel=mid_channel, kernel_size=1, stride=1, groups=1),
        )
        if self.stride > 1:
            # 论文中f(g,3,d) 中的左半部分
            self.shortcut = nn.Sequential(
                # 3x3 DWConv
                ConvBNReLU(in_channel=in_channel, out_channel=in_channel, kernel_size=1, stride=1, groups=1),
            )
        self.channel_shuffle = ChannelShuffle(groups=groups)

    def forward(self, x):
        if self.stride > 1:
            x1 = self.bottleneck(x)
            x2 = self.shortcut(x)
        else:
            # channel split
            x1 = self.first_split(x)
            x2 = self.second_split(x)
            x1 = self.bottleneck(x1)
            out = torch.cat([x1, x2], dim=1)
        out = self.channel_shuffle(out)
        return out

class ShuffleNetV2(nn.Module):
    def __init__(self, planes, layers, groups, num_classes=1000):
        super(ShuffleNetV2, self).__init__()
        self.groups = groups

        self.stage1 = nn.Sequential(
            ConvBNReLU(in_channel=3, out_channel=24, kernel_size=3, stride=2, groups=1),
            nn.MaxPool2d(kernel_size=3, stride=2, padding=1),
        )
        self.stage2 = self._make_layers(in_channel=planes[0], out_channel=planes[1], block_num=layers[0], self.stage2=False)
        self.stage3 = self._make_layers(in_channel=planes[1], out_channel=planes[2], block_num=layers[1], self.stage2=False)
        self.stage4 = self._make_layers(in_channel=planes[2], out_channel=planes[3], block_num=layers[2], self.stage2=False)

        self.conv5 = ConvBNReLU(in_channel=planes[3], out_channel=planes[3], kernel_size=1, stride=1, groups=1)
        self.globalpool = nn.AdaptiveAvgPool2d([1, 1])
        self.fc = nn.Sequential(
            nn.Dropout(p=0.2),
            nn.Linear(in_features=planes[3], out_features=num_classes)
        )

        for m in self.modules():
            if isinstance(m, nn.Conv2d):
                nn.init.kaiming_normal_(m.weight)
            elif isinstance(m, nn.BatchNorm2d) or isinstance(m, nn.Linear):
                nn.init.constant_(m.weight, 1)
                nn.init.constant_(m.bias, 0)

        def _make_layers(self, in_channel, out_channel, block_num, is_stage2):
            layers.append(ShuffleNetUnits(in_channel=in_channel, out_channel=out_channel, stride=2, groups=self.groups))
            for i in range(1, block_num):
                ShuffleNetUnits(in_channel=out_channel, out_channel=out_channel, stride=1, groups=self.groups)
            return nn.Sequential(*layers)

    def forward(self, x):
        x = self.stage1(x)
        x = self.stage2(x)
        x = self.stage3(x)
        x = self.stage4(x)
        x = self.conv5(x)
        x = self.globalpool(x)
        x = torch.flatten(x, start_dim=1)
        x = self.fc(x)
        return x

    def shufflenet_v2_x1_5(**kwargs):
        planes = [176, 352, 704, 1024]
        layers = [4, 8, 4]
        model = ShuffleNetV2(planes=planes, layers=layers, groups=1, **kwargs)
        return model

In [8]: inputs = torch.randn(1, 3, 224, 224)
model_shufflenet = shufflenet_v2_x1_5(num_classes=10)
out = model_shufflenet(inputs)
print(out.shape)

torch.Size([1, 10])
```

## 绘图函数

plot\_model\_losses 函数可以将多个模型的loss训练过程画成一张图，便于展示结果。

```
In [9]: def plot_model_losses(train_losses, val_losses, names):
    nets = len(names)
    cols = 3

    rows = (nets + cols - 1) // cols # 计算行数, 确保布局合理

    # 创建一个符合制布局的图形
    fig, axes = plt.subplots(rows, cols, figsize=(cols * 6, rows * 5))

    # if rows == 1:
    #     axes = [axes] # 如果只有一行, 确保 axes 是一个列表

    for i in range(nets):
        row = i // cols
        col = i % cols
        ax = axes[row][col] if rows > 1 else axes[col]

        # 确保 train_losses[i] 和 val_losses[i] 是一维数组或列表
        ax.plot(train_losses[i], label='Train Loss')
        ax.plot(val_losses[i], label='Test Loss')
        ax.set_title(f'Model {names[i]} Losses')
        ax.set_xlabel('Epoch')
        ax.set_ylabel('Loss')
        ax.legend()

    plt.tight_layout()

    # 关闭多余的子图
    for j in range(1 + rows * cols):
        row = j // cols
        col = j % cols
        ax = axes[row][col] if rows > 1 else axes[col]
        ax.axis('off')

    plt.show()

plot_val_acc 函数可以接受模型的验证准确率、模型名称、绘图样式，然后将所有模型的验证集准确率进行对比
```

```
In [10]: def plot_val_acc(val_acc_list, model_names, styles=None, figsize=(15, 5), title='Model Accuracy', y_labels=None):
    if styles is None:
        styles = ['-', '--', '-', '-', '-', '-', '-', '-', '-'] * (len(val_acc_list) // 4 + 1)

    plt.figure(figsize=figsize)
    for i, val_acc in enumerate(val_acc_list):
        plt.plot(val_acc, linestyle=styles[i % len(styles)], label=model_names[i])

    plt.title(title)
    plt.xlabel('Epoch')
    plt.ylabel(val_acc)
    axes = plt.gca()
    axes.set_ylim(0, 1)
    plt.show()
```

## 模型训练函数

```
In [11]: def train_model(model, train_loader, test_loader, epochs=20, device="cpu", model_name=""):
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters())

    train_acc = []
    val_acc = []
    train_loss = []
    val_loss = []

    model.to(device)

    while True:
        folder_path = f'./weight_{model_name} + '_' + ''.join(random.choices(string.ascii_lowercase, k=5))
        if not os.path.exists(folder_path):
            os.makedirs(folder_path)
            break

        print(f'weights are saved to {folder_path} ")

        for epoch in range(epochs):
            # Training phase
            model.train()
            correct = 0
            total = 0
            running_loss = 0.0
            for images, labels in train_loader:
                images, labels = images.to(device), labels.to(device)

                optimizer.zero_grad()
                outputs = model(images)
                loss = criterion(outputs, labels)
                loss.backward()
                optimizer.step()

                running_loss += loss.item()
                _, predicted = torch.max(outputs.data, 1)
                total += labels.size(0)
                correct += (predicted == labels).sum().item()

            train_loss.append(running_loss / len(train_loader))
            train_acc.append(correct / total)

            # Validation phase
            model.eval()
            correct = 0
            total = 0
            running_loss = 0.0
            with torch.no_grad():
                for images, labels in test_loader:
                    images, labels = images.to(device), labels.to(device)
                    outputs = model(images)
                    loss = criterion(outputs, labels)

                    running_loss += loss.item()
                    _, predicted = torch.max(outputs.data, 1)
                    total += labels.size(0)
                    correct += (predicted == labels).sum().item()

            val_loss.append(running_loss / len(test_loader))
            val_acc.append(correct / total)

        # 保存模型
        model_file = os.path.join(folder_path, f"epoch_{epoch+1}.pth")
        torch.save(model.state_dict(), model_file)
        print(f"Saved model at epoch {epoch+1} to {model_file}")

        print(f"Epoch {epoch+1} of {epochs}, Train Loss: {train_loss[-1:4f]}, Train Accuracy: {train_val_acc[-1:4f]}, Train Loss: {val_loss[-1:4f]}, Train Accuracy: {val_acc[-1:4f]}")

    return train_loss, train_acc, val_loss, val_acc
```

## 数据集导入

```
In [12]: # 强力的数据增强组合, 只应用于训练集
train_transform = transforms.Compose([
    transforms.Resize(size=224),
    transforms.RandomHorizontalFlip(), # 随机水平翻转, 保持不变
    transforms.RandomRotation(10), # 随机旋转角度减小, 从15度减小到5度
    transforms.RandomResizedCrop(224, scale=(0.85, 1.0)), # 随机裁剪
    transforms.ColorJitter(brightness=0.1, contrast=0.1, saturation=0.1, hue=0.1), # 随机调整颜色偏色
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)) # 归一化
])

# 仅进行标准化的转换, 应用于测试集
test_transform = transforms.Compose([
    transforms.Resize(size=224),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)) # 归一化
])

# 加载CIFAR-10数据集
train_dataset = datasets.CIFAR10(root='./data', train=True, download=True, transform=train_transform)
test_dataset = datasets.CIFAR10(root='./data', train=False, download=True, transform=test_transform)

train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)

Downloading https://www.cs.toronto.edu/~kriz/cifar-100-python.tar.gz to ./data/cifar-100-python.tar.gz
100%|#####| 51698071/51698071 [06:03<00:00, 49241583.26it/s]
Extracting ./data/cifar-100-python.tar.gz to ./data
Files already downloaded and verified
```

```
In [13]: fig, axs = plt.subplots(1,5,figsize=(10,10)) #建立子图
for i in range(5):
    num = random.randint(0,len(data)-1) # 首先选取随机数, 随机选取五次
    # 提取数据并返回的图像数据, make_grid函数将任意格式的图像的通道数升为3, 而不改变图像原始的数据
    # 而展示图像用的imshow函数最常见的输入格式也是3通道
    npimg = torchvision.utils.make_grid(data[num][0]).numpy()
    plt.imshow(npimg) # 提取数据
    # 将图像(x, y, weight, height)转成列(weight, height, num, 3), 并放入imshow函数中读取
    axs[i].set_title(npimgLabel) # 给每个子图加上标签
    axs[i].axis('off') # 隐藏每个子图的坐标轴
```

```
In [14]: plotsample(train_dataset)

8      6      7      5      6
```



```
In [15]: plotsample(test_dataset)

3      6      1      0      7
```



## 训练

```
In [16]: model_name_list = ['alex', 'Mobilenetv3', 'shufflenet']
train_loss_list, train_acc_list, val_loss_list, val_acc_list = [], [], [], []
```

```
In [17]: train_loss, train_acc, val_loss, val_acc = train_model(model=model_AlexNet, train_loader=train_loader, test_loader=test_loader, epochs=20, device="cpu", model_name="alex")
train_loss_list.append(train_loss)
train_acc_list.append(train_acc)
val_loss_list.append(val_loss)
val_acc_list.append(val_acc)
```

```
weights are saved to ./weight_alex_zdujejk25c
Saved model at epoch 1 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_1.pth
Epoch 1/15, Train Loss: 1.8127, Train Accuracy: 0.3208, Validation Loss: 1.5210, Validation Accuracy: 0.4327
Saved model at epoch 2 to ./weight_alex_zdujejk25c/epoch_2.pth
Epoch 2/15, Train Loss: 1.5230, Train Accuracy: 0.4433, Validation Loss: 1.3262, Validation Accuracy: 0.6069
Saved model at epoch 3 to ./weight_alex_zdujejk25c/epoch_3.pth
Epoch 3/15, Train Loss: 0.7635, Train Accuracy: 0.4925, Validation Loss: 1.2323, Validation Accuracy: 0.6251
Saved model at epoch 4 to ./weight_alex_zdujejk25c/epoch_4.pth
Epoch 4/15, Train Loss: 1.3386, Train Accuracy: 0.5251, Validation Loss: 1.1336, Validation Accuracy: 0.6690
Saved model at epoch 5 to ./weight_alex_zdujejk25c/epoch_5.pth
Epoch 5/15, Train Loss: 1.2715, Train Accuracy: 0.5472, Validation Loss: 1.1187, Validation Accuracy: 0.6647
Saved model at epoch 6 to ./weight_alex_zdujejk25c/epoch_6.pth
Epoch 6/15, Train Loss: 1.2211, Train Accuracy: 0.5689, Validation Loss: 1.0880, Validation Accuracy: 0.6207
Saved model at epoch 7 to ./weight_alex_zdujejk25c/epoch_7.pth
Epoch 7/15, Train Loss: 1.1293, Train Accuracy: 0.5748, Validation Loss: 1.0370, Validation Accuracy: 0.6381
Saved model at epoch 8 to ./weight_alex_zdujejk25c/epoch_8.pth
Epoch 8/15, Train Loss: 1.1598, Train Accuracy: 0.5954, Validation Loss: 1.0525, Validation Accuracy: 0.6286
Saved model at epoch 9 to ./weight_alex_zdujejk25c/epoch_9.pth
Epoch 9/15, Train Loss: 1.1293, Train Accuracy: 0.6054, Validation Loss: 0.9614, Validation Accuracy: 0.6692
Saved model at epoch 10 to ./weight_alex_zdujejk25c/epoch_10.pth
Epoch 10/15, Train Loss: 1.0973, Train Accuracy: 0.6175, Validation Loss: 0.9591, Validation Accuracy: 0.6657
Saved model at epoch 11 to ./weight_alex_zdujejk25c/epoch_11.pth
Epoch 11/15, Train Loss: 1.0824, Train Accuracy: 0.6214, Validation Loss: 0.9518, Validation Accuracy: 0.6715
Saved model at epoch 12 to ./weight_alex_zdujejk25c/epoch_12.pth
Epoch 12/15, Train Loss: 1.0629, Train Accuracy: 0.6295, Validation Loss: 0.9080, Validation Accuracy: 0.6542
Saved model at epoch 13 to ./weight_alex_zdujejk25c/epoch_13.pth
Epoch 13/15, Train Loss: 1.0421, Train Accuracy: 0.6373, Validation Loss: 0.9384, Validation Accuracy: 0.6824
Saved model at epoch 14 to ./weight_alex_zdujejk25c/epoch_14.pth
Epoch 14/15, Train Loss: 1.0254, Train Accuracy: 0.6451, Validation Loss: 0.9043, Validation Accuracy: 0.6895
Saved model at epoch 15 to ./weight_alex_zdujejk25c/epoch_15.pth
Epoch 15/15, Train Loss: 1.0169, Train Accuracy: 0.6469, Validation Loss: 0.9034, Validation Accuracy: 0.6826
```

```
In [18]: train_loss, train_acc, val_loss, val_acc = train_model(model=model_Mobilenetv3, train_loader=train_loader, test_loader=test_loader, epochs=20, device="cpu", model_name="mobilenetv3")
train_loss_list.append(train_loss)
train_acc_list.append(train_acc)
val_loss_list.append(val_loss)
val_acc_list.append(val_acc)
```

```
weights are saved to ./weight_Mobilenetv3_2qxxfmfp8
Saved model at epoch 1 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_1.pth
Epoch 1/15, Train Loss: 1.7753, Train Accuracy: 0.3233, Validation Loss: 1.3833, Validation Accuracy: 0.4876
Saved model at epoch 2 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_2.pth
Epoch 2/15, Train Loss: 1.3119, Train Accuracy: 0.4776, Validation Loss: 1.1854, Validation Accuracy: 0.5834
Saved model at epoch 3 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_3.pth
Epoch 3/15, Train Loss: 0.7933, Train Accuracy: 0.6572, Validation Loss: 1.0460, Validation Accuracy: 0.6235
Saved model at epoch 4 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_4.pth
Epoch 4/15, Train Loss: 0.9733, Train Accuracy: 0.6572, Validation Loss: 0.8541, Validation Accuracy: 0.6069
Saved model at epoch 5 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_5.pth
Epoch 5/15, Train Loss: 0.8226, Train Accuracy: 0.7094, Validation Loss: 0.7359, Validation Accuracy: 0.7421
Saved model at epoch 6 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_6.pth
Epoch 6/15, Train Loss: 0.7740, Train Accuracy: 0.7314, Validation Loss: 0.6752, Validation Accuracy: 0.7687
Saved model at epoch 7 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_7.pth
Epoch 7/15, Train Loss: 0.6160, Train Accuracy: 0.7712, Validation Loss: 0.6909, Validation Accuracy: 0.7593
Saved model at epoch 8 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_8.pth
Epoch 8/15, Train Loss: 0.7655, Train Accuracy: 0.8049, Validation Loss: 0.6534, Validation Accuracy: 0.6985
Saved model at epoch 9 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_9.pth
Epoch 9/15, Train Loss: 0.6526, Train Accuracy: 0.7768, Validation Loss: 0.6795, Validation Accuracy: 0.7975
Saved model at epoch 10 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_10.pth
Epoch 10/15, Train Loss: 0.6269, Train Accuracy: 0.7946, Validation Loss: 0.5647, Validation Accuracy: 0.8092
Saved model at epoch 11 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_11.pth
Epoch 11/15, Train Loss: 0.5904, Train Accuracy: 0.7946, Validation Loss: 0.6045, Validation Accuracy: 0.7976
Saved model at epoch 12 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_12.pth
Epoch 12/15, Train Loss: 0.5607, Train Accuracy: 0.8065, Validation Loss: 0.5082, Validation Accuracy: 0.8299
Saved model at epoch 13 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_13.pth
Epoch 13/15, Train Loss: 0.5082, Train Accuracy: 0.8062, Validation Loss: 0.5450, Validation Accuracy: 0.8206
Saved model at epoch 14 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_14.pth
Epoch 14/15, Train Loss: 0.4691, Train Accuracy: 0.8376, Validation Loss: 0.4634, Validation Accuracy: 0.8399
Saved model at epoch 15 to ./weight_Mobilenetv3_2qxxfmfp8/epoch_15.pth
Epoch 15/15, Train Loss: 0.4358, Train Accuracy: 0.8495, Validation Loss: 0.5034, Validation Accuracy: 0.8286
```

```
In [19]: train_loss, train_acc, val_loss, val_acc = train_model(model=model_shufflenet, train_loader=train_loader, test_loader=test_loader, epochs=20, device="cpu", model_name="shufflenet")
train_loss_list.append(train_loss)
train_acc_list.append(train_acc)
val_loss_list.append(val_loss)
val_acc_list.append(val_acc)
```



weights are saved to ./weight\_shuffleNet\_t5vlyirho1  
Saved model at epoch 1 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_1.pth  
Epoch 2/15, Train Loss: 1.5628, Train Accuracy: 0.4185, Validation Loss: 1.2502, Validation Accuracy: 0.5494  
Saved model at epoch 2 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_2.pth  
Epoch 3/15, Train Loss: 1.1423, Train Accuracy: 0.5922, Validation Loss: 0.9566, Validation Accuracy: 0.6592  
Saved model at epoch 3 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_3.pth  
Epoch 4/15, Train Loss: 0.9214, Train Accuracy: 0.6751, Validation Loss: 0.8215, Validation Accuracy: 0.7066  
Saved model at epoch 4 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_4.pth  
Epoch 5/15, Train Loss: 0.7834, Train Accuracy: 0.7252, Validation Loss: 0.7034, Validation Accuracy: 0.7578  
Saved model at epoch 5 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_5.pth  
Epoch 6/15, Train Loss: 0.6245, Train Accuracy: 0.7832, Validation Loss: 0.5982, Validation Accuracy: 0.7959  
Saved model at epoch 6 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_6.pth  
Epoch 7/15, Train Loss: 0.5661, Train Accuracy: 0.8032, Validation Loss: 0.5492, Validation Accuracy: 0.8128  
Saved model at epoch 7 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_8.pth  
Epoch 8/15, Train Loss: 0.5237, Train Accuracy: 0.8177, Validation Loss: 0.5490, Validation Accuracy: 0.8187  
Saved model at epoch 8 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_9.pth  
Epoch 9/15, Train Loss: 0.4886, Train Accuracy: 0.8314, Validation Loss: 0.4955, Validation Accuracy: 0.8297  
Saved model at epoch 9 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_10.pth  
Epoch 10/15, Train Loss: 0.4544, Train Accuracy: 0.8433, Validation Loss: 0.4809, Validation Accuracy: 0.8347  
Saved model at epoch 10 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_11.pth  
Epoch 11/15, Train Loss: 0.4286, Train Accuracy: 0.8516, Validation Loss: 0.4722, Validation Accuracy: 0.8387  
Saved model at epoch 11 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_12.pth  
Epoch 12/15, Train Loss: 0.4017, Train Accuracy: 0.8601, Validation Loss: 0.4565, Validation Accuracy: 0.8483  
Saved model at epoch 12 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_13.pth  
Epoch 13/15, Train Loss: 0.3805, Train Accuracy: 0.8668, Validation Loss: 0.4584, Validation Accuracy: 0.8455  
Saved model at epoch 13 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_14.pth  
Epoch 14/15, Train Loss: 0.3676, Train Accuracy: 0.8729, Validation Loss: 0.4552, Validation Accuracy: 0.8483  
Saved model at epoch 14 to ./weight\_shuffleNet\_t5vlyirho1/epoch\_15.pth  
Epoch 15/15, Train Loss: 0.3443, Train Accuracy: 0.8808, Validation Loss: 0.4140, Validation Accuracy: 0.8632

In [20]: plot\_model\_losses(train\_loss\_list, val\_loss\_list, model\_name\_list)



In [21]: plot\_val\_acc(val\_acc\_list, model\_name\_list, styles=None, figsize=(15, 5), title='Model Accuracy', )



In [ ]: plot\_val\_acc(val\_acc\_list, model\_name\_list, styles=None, figsize=(15, 5), title='Model Accuracy', )