

消融实验

Knotebook是在清华大学在24Fall江苏大学人工智能专业实践的小组作业。基于cifar10数据集的模型消融实验。主要满足有目的的移动或替换模型中的某些组件或功能。研究这些组件对模型性能的影响，从而理解模型的关键特性和有效性。

本次实验探究了以下五种因素对模型的影响：

- convolution-subsampling pairs
- featuremap num
- dense layers
- dropout
- Augmentation

导入本次实验的库

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader

In [2]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
device

Out[2]: device(type='cuda')
```

导入数据集

```
In [3]: # Load CIFAR-10
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

train_dataset = datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
test_dataset = datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)

train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)

Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to ./data/cifar-10-python.tar.gz
100%|#####| 170480871/170480871 [00:03:00.00, 49038513.09it/s]
Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified
```

创建绘图函数

```
plot_model_losses 函数可以将多个模型的loss训练过程绘图在一张图，便于展示结果。

In [4]: def plot_model_losses(train_losses, val_losses, names):
    nets = len(names)
    cols = 3

    rows = (nets + cols - 1) // cols # 计算行数，确保布局合理

    # 创建一个符合网格布局的图形
    fig, axes = plt.subplots(rows, cols, figsize=(cols * 6, rows * 5))

    # 如果 rows == 1:
    #     axes = [axes] # 如果只有一行，确保 axes 是一个列表

    for i in range(nets):
        row = i // cols
        col = i % cols
        ax = axes[row][col] if rows > 1 else axes[col]

        # 确保 train_losses[i] 和 val_losses[i] 是一维数组或列表
        ax.plot(train_losses[i], label='Train Loss')
        ax.plot(val_losses[i], label='Test Loss')
        ax.set_title(f'Model {names[i]} Losses')
        ax.set_xlabel('Epoch')
        ax.set_ylabel('Loss')
        ax.legend()

    plt.tight_layout()

    # 关闭空的子图
    for j in range(1 + 1, rows + cols):
        row = j // cols
        col = j % cols
        ax = axes[row][col] if rows > 1 else axes[col]
        ax.axis('off')

    plt.show()

plot_val_acc函数可以接受模型的验证准确率、模型名称、绘图样式，然后将所有模型的验证准确率进行对比

In [5]: def plot_val_acc(val_acc_list, model_names, styles=None, figsize=(15, 5), title='Model Accuracy', y_styles = ['-', '--', '-', '-', '-', '-', '-', '-', '-'], x_styles = ['-', '-', '-', '-', '-', '-', '-', '-', '-']):
    if styles is None:
        styles = [ '-', '-', '-', '-', '-', '-', '-', '-', '-'] * (len(val_acc_list) // 4 + 1)

    plt.figure(figsize=figsize)
    for i, val_acc in enumerate(val_acc_list):
        plt.plot(val_acc, linestyle=styles[i % len(styles)], label=model_names[i])

    plt.title(title)
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.legend(loc='upper left')
    axes = plt.gca().xaxis
    axes.set_ylim(0, 1)
    plt.show()
```

模型训练函数定义

定义模型训练的函数。返回 train_loss, train_acc, val_loss, val_acc
模型的optimizer使用Adam, criterion使用交叉熵损失

```
In [6]: def train_model(model, train_loader, test_loader, epochs=20, device="cpu"):
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=0.001)

    train_acc = []
    val_acc = []
    train_loss = []
    val_loss = []

    for epoch in range(epochs):
        # training phase
        model.train()
        correct = 0
        total = 0
        running_loss = 0.0
        for images, labels in train_loader:
            images, labels = images.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.item()
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

        train_loss.append(running_loss / len(train_loader))
        val_acc.append(correct / total)

        # Validation phase
        model.eval()
        correct = 0
        total = 0
        running_loss = 0.0
        with torch.no_grad():
            for images, labels in test_loader:
                images, labels = images.to(device), labels.to(device)
                outputs = model(images)
                loss = criterion(outputs, labels)

                running_loss += loss.item()
                _, predicted = torch.max(outputs.data, 1)
                total += labels.size(0)
                correct += (predicted == labels).sum().item()

        val_loss.append(running_loss / len(test_loader))
        val_acc.append(correct / total)

    if epoch % 2 == 0 or epoch + 1 == epochs:
        print(f'Epoch {epoch+1}/{epochs}, Train Loss: {train_loss[-1]:.4f}, Train Accuracy: {train_acc[-1]:.4f}, Val Loss: {val_loss[-1]:.4f}, Val Accuracy: {val_acc[-1]:.4f}')

    return train_loss, train_acc, val_loss, val_acc
```

比较试验

1.探究 convolution-subsampling pairs 对模型影响

建立一个简单的模型，探究不同的 convolution-subsampling pairs 对数对模型的影响，我们比较了四种参数的组合，通过堆叠（32x32的卷积层+2x2的池化层）：

- 32CP2
- 32CP2-32CP2
- 32CP2-32CP2-32CP2
- 4.32CP2-32CP2-32CP2-32CP2

模型定义

```
In [7]: class SimpleCNN(nn.Module):
    def __init__(self, conv_subsample_pairs):
        super(SimpleCNN, self).__init__()
        layers = []
        in_channels = 3
        out_channels = 32

        for i in range(conv_subsample_pairs):
            layers.append(nn.Conv2d(in_channels, out_channels, kernel_size=5, stride=1, padding=2))
            layers.append(nn.ReLU())
            layers.append(nn.MaxPool2d(kernel_size=2, stride=2))
            in_channels = out_channels
            out_channels = 16

        out_channels = 16
        layers.append(nn.Flatten())
        # Compute the size of the input to the linear layer (taking into account the pooling)
        linear_input_size = out_channels * (32 // (2 ** conv_subsample_pairs)) ** 2
        layers.append(nn.Linear(linear_input_size, 10))

        self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)
```

实验过程

```
In [8]: pairs_list = [1, 2, 3, 4]
names = [f'pairs {pairs}' for pairs in pairs_list]
styles = ['-', '-', '-', '-']
train_loss_list = []
train_acc_list = []
val_loss_list = []
val_acc_list = []

for pairs in pairs_list:
    print(f'Training with {pairs} convolution-subsampling pairs...')
    model = SimpleCNN(conv_subsample_pairs=pairs).to(device)
    train_loss, train_acc, val_loss, val_acc = train_model(model, train_loader, test_loader, epochs=20)
    train_loss_list.append(train_loss)
    train_acc_list.append(train_acc)
    val_loss_list.append(val_loss)
    val_acc_list.append(val_acc)

Training with 1 convolution-subsampling pairs...
Epoch 1/20, Train Loss: 1.3818, Train Accuracy: 0.5179, Validation Loss: 1.1692, Validation Accuracy: 0.5905
Epoch 3/20, Train Loss: 0.9799, Train Accuracy: 0.6658, Validation Loss: 1.0447, Validation Accuracy: 0.6421
Epoch 5/20, Train Loss: 0.8724, Train Accuracy: 0.7814, Validation Loss: 1.0099, Validation Accuracy: 0.6563
Epoch 7/20, Train Loss: 0.8006, Train Accuracy: 0.7235, Validation Loss: 1.0607, Validation Accuracy: 0.6448
Epoch 9/20, Train Loss: 0.7351, Train Accuracy: 0.7486, Validation Loss: 1.0446, Validation Accuracy: 0.6540
Epoch 11/20, Train Loss: 0.6897, Train Accuracy: 0.7621, Validation Loss: 1.0676, Validation Accuracy: 0.6510
Epoch 13/20, Train Loss: 0.6447, Train Accuracy: 0.7773, Validation Loss: 1.0926, Validation Accuracy: 0.6588
Epoch 15/20, Train Loss: 0.5692, Train Accuracy: 0.7983, Validation Loss: 1.2048, Validation Accuracy: 0.6480
Epoch 17/20, Train Loss: 0.4654, Train Accuracy: 0.8038, Validation Loss: 1.2488, Validation Accuracy: 0.6583
Epoch 19/20, Train Loss: 0.3714, Train Accuracy: 0.8128, Validation Loss: 1.2488, Validation Accuracy: 0.6480
Epoch 20/20, Train Loss: 0.5244, Train Accuracy: 0.8178, Validation Loss: 1.2282, Validation Accuracy: 0.6375
Training with 2 convolution-subsampling pairs...
Epoch 1/20, Train Loss: 1.1429, Validation Loss: 1.0762, Validation Accuracy: 0.6224
Epoch 3/20, Train Loss: 0.8391, Train Accuracy: 0.8383, Validation Loss: 0.8692, Validation Accuracy: 0.7039
Epoch 5/20, Train Loss: 0.6852, Train Accuracy: 0.7825, Validation Loss: 0.8678, Validation Accuracy: 0.7229
Epoch 7/20, Train Loss: 0.5915, Train Accuracy: 0.7945, Validation Loss: 0.8225, Validation Accuracy: 0.7285
Epoch 9/20, Train Loss: 0.5155, Train Accuracy: 0.8287, Validation Loss: 0.8593, Validation Accuracy: 0.7233
Epoch 11/20, Train Loss: 0.4510, Train Accuracy: 0.8409, Validation Loss: 0.9100, Validation Accuracy: 0.7219
Epoch 13/20, Train Loss: 0.3986, Train Accuracy: 0.8383, Validation Loss: 0.9556, Validation Accuracy: 0.7244
Epoch 15/20, Train Loss: 0.3563, Train Accuracy: 0.8747, Validation Loss: 1.0499, Validation Accuracy: 0.7153
Epoch 17/20, Train Loss: 0.3153, Train Accuracy: 0.8888, Validation Loss: 1.1683, Validation Accuracy: 0.7088
Epoch 19/20, Train Loss: 0.2778, Train Accuracy: 0.9481, Validation Loss: 1.2234, Validation Accuracy: 0.7108
Epoch 20/20, Train Loss: 0.1870, Train Accuracy: 0.9088, Validation Loss: 1.2706, Validation Accuracy: 0.7098
Training with 3 convolution-subsampling pairs...
Epoch 1/20, Train Loss: 1.4008, Train Accuracy: 0.4926, Validation Loss: 1.0932, Validation Accuracy: 0.6059
Epoch 3/20, Train Loss: 0.8093, Train Accuracy: 0.7189, Validation Loss: 0.8617, Validation Accuracy: 0.6967
Epoch 5/20, Train Loss: 0.6043, Train Accuracy: 0.7983, Validation Loss: 0.8218, Validation Accuracy: 0.7283
Epoch 7/20, Train Loss: 0.4654, Train Accuracy: 0.8363, Validation Loss: 0.8296, Validation Accuracy: 0.7317
Epoch 9/20, Train Loss: 0.3545, Train Accuracy: 0.8761, Validation Loss: 0.9263, Validation Accuracy: 0.7175
Epoch 11/20, Train Loss: 0.2699, Train Accuracy: 0.9045, Validation Loss: 1.0519, Validation Accuracy: 0.7263
Epoch 13/20, Train Loss: 0.2131, Train Accuracy: 0.9253, Validation Loss: 1.1839, Validation Accuracy: 0.7340
Epoch 15/20, Train Loss: 0.1806, Train Accuracy: 0.9352, Validation Loss: 1.3708, Validation Accuracy: 0.7261
Epoch 17/20, Train Loss: 0.1460, Train Accuracy: 0.9481, Validation Loss: 1.5741, Validation Accuracy: 0.7178
Epoch 19/20, Train Loss: 0.1352, Train Accuracy: 0.9521, Validation Loss: 1.6596, Validation Accuracy: 0.7236
Epoch 20/20, Train Loss: 0.1231, Train Accuracy: 0.9559, Validation Loss: 1.8823, Validation Accuracy: 0.7208
Training with 4 convolution-subsampling pairs...
Epoch 1/20, Train Loss: 1.3571, Validation Loss: 1.1034, Validation Accuracy: 0.6057
Epoch 3/20, Train Loss: 0.7710, Train Accuracy: 0.7383, Validation Loss: 0.8385, Validation Accuracy: 0.7119
Epoch 5/20, Train Loss: 0.5395, Train Accuracy: 0.8112, Validation Loss: 0.8040, Validation Accuracy: 0.7283
Epoch 7/20, Train Loss: 0.3727, Train Accuracy: 0.8675, Validation Loss: 0.8422, Validation Accuracy: 0.7410
Epoch 9/20, Train Loss: 0.2436, Train Accuracy: 0.9137, Validation Loss: 1.0891, Validation Accuracy: 0.7207
Epoch 11/20, Train Loss: 0.1497, Train Accuracy: 0.9313, Validation Loss: 1.4511, Validation Accuracy: 0.7314
Epoch 13/20, Train Loss: 0.1387, Train Accuracy: 0.9510, Validation Loss: 1.4151, Validation Accuracy: 0.7348
Epoch 15/20, Train Loss: 0.1290, Train Accuracy: 0.9556, Validation Loss: 1.5842, Validation Accuracy: 0.7328
Epoch 17/20, Train Loss: 0.1369, Train Accuracy: 0.9618, Validation Loss: 1.6246, Validation Accuracy: 0.7317
Epoch 19/20, Train Loss: 0.1137, Train Accuracy: 0.9619, Validation Loss: 1.8886, Validation Accuracy: 0.7248
Epoch 20/20, Train Loss: 0.1139, Train Accuracy: 0.9618, Validation Loss: 1.7957, Validation Accuracy: 0.7293
```

实验结果

```
In [9]: plot_val_acc(val_acc_list, names, styles=styles, title='Effect of Convolution-Subsampling Pairs on CIFAR-10')

Effect of Convolution-Subsampling Pairs on CIFAR-10
```

```
In [10]: plot_model_losses(train_loss_list, val_loss_list, names)

Model 1 pairs Losses      Model 2 pairs Losses      Model 3 pairs Losses
Accuracy vs Epoch
```

可以看到，我们使用三次堆叠的结果就可以在当前的状态下得到很好的结果。

2.探究 feature map nums 对模型影响

探究特征图通深度的影响，组合如下：

- 16-32-48, 32-48-64, 64-96-128

模型定义

```
In [11]: class SimpleCNN(nn.Module):
    def __init__(self, channels=None):
        super(SimpleCNN, self).__init__()
        layers = []
        in_channels = 3

        # 设置默认的快速通道配置
        if channels_config is None:
            channels_config = [16, 32, 48] # 默认配置: 16 -> 32 -> 48

        # 根据conv_subsample_pairs和channels配置设置卷积层
        for i in range(3):
            out_channels = channels_config[i]
            layers.append(nn.Conv2d(in_channels, out_channels, kernel_size=5, stride=1, padding=2))
            layers.append(nn.ReLU())
            layers.append(nn.MaxPool2d(kernel_size=2, stride=2))
            in_channels = out_channels

        layers.append(nn.Flatten())

        linear_input_size = in_channels * 16
        layers.append(nn.Linear(linear_input_size, 10))

        self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)
```

实验过程

```
In [12]: # 实验过程
channel_configs = [
    [16, 32, 48], # 组合1
    [32, 48, 64], # 组合2
    [64, 96, 128] # 组合3
]
names = [f'Channels: {cfg}' for cfg in channel_configs]
styles = ['-', '-', '-', '-']

train_loss_list = []
train_acc_list = []
val_loss_list = []
val_acc_list = []

In [13]: # 进行实验
for channel_config in channel_configs:
    print(f'Training with channel config {channel_config}...')
    model = SimpleCNN(channels_config=channel_config).to(device) # 假设使用预置卷积
    train_loss, train_acc, val_loss, val_acc = train_model(model, train_loader, test_loader, epochs=20)
    train_loss_list.append(train_loss)
    train_acc_list.append(train_acc)
    val_loss_list.append(val_loss)
    val_acc_list.append(val_acc)

Training with channel config [16, 32, 48]...
Epoch 1/10, Train Loss: 1.4409, Train Accuracy: 0.4772, Validation Loss: 1.2008, Validation Accuracy: 0.5697
Epoch 3/10, Train Loss: 0.9638, Train Accuracy: 0.6826, Validation Loss: 0.8988, Validation Accuracy: 0.6854
Epoch 5/10, Train Loss: 0.7691, Train Accuracy: 0.7491, Validation Loss: 0.8335, Validation Accuracy: 0.7086
Epoch 7/10, Train Loss: 0.6109, Train Accuracy: 0.7886, Validation Loss: 0.8067, Validation Accuracy: 0.7194
Epoch 9/10, Train Loss: 0.5301, Train Accuracy: 0.8143, Validation Loss: 0.8605, Validation Accuracy: 0.7263
Epoch 10/10, Train Loss: 0.4927, Train Accuracy: 0.8294, Validation Loss: 0.8607, Validation Accuracy: 0.7194
Training with channel config [32, 48, 64]...
Epoch 1/10, Train Loss: 1.3877, Train Accuracy: 0.4993, Validation Loss: 1.1196, Validation Accuracy: 0.6078
Epoch 3/10, Train Loss: 0.8048, Train Accuracy: 0.7203, Validation Loss: 0.8600, Validation Accuracy: 0.7054
Epoch 5/10, Train Loss: 0.6145, Train Accuracy: 0.7863, Validation Loss: 0.7907, Validation Accuracy: 0.7362
Epoch 7/10, Train Loss: 0.4815, Train Accuracy: 0.8386, Validation Loss: 0.8165, Validation Accuracy: 0.7400
Epoch 9/10, Train Loss: 0.3372, Train Accuracy: 0.8828, Validation Loss: 0.9579, Validation Accuracy: 0.7237
Epoch 10/10, Train Loss: 0.3173, Train Accuracy: 0.8928, Validation Loss: 0.9579, Validation Accuracy: 0.7237
Training with channel config [64, 96, 128]...
Epoch 1/10, Train Loss: 1.2979, Train Accuracy: 0.5336, Validation Loss: 0.9865, Validation Accuracy: 0.6550
Epoch 3/10, Train Loss: 0.7966, Train Accuracy: 0.7576, Validation Loss: 0.7938, Validation Accuracy: 0.7091
Epoch 5/10, Train Loss: 0.4601, Train Accuracy: 0.8391, Validation Loss: 0.7734, Validation Accuracy: 0.7505
Epoch 7/10, Train Loss: 0.2842, Train Accuracy: 0.9015, Validation Loss: 0.8819, Validation Accuracy: 0.7491
Epoch 9/10, Train Loss: 0.1913, Train Accuracy: 0.9334, Validation Loss: 1.0969, Validation Accuracy: 0.7298
Epoch 10/10, Train Loss: 0.1584, Train Accuracy: 0.9432, Validation Loss: 1.2290, Validation Accuracy: 0.7437
```

实验结果

```
In [14]: # 绘制训练验证曲线
plot_val_acc(val_acc_list, names, styles=styles, title='Effect of Different Channel Configurations')

Effect of Different Channel Configurations on CIFAR-10
```

```
In [15]: # 绘制训练和验证损失曲线
plot_model_losses(train_loss_list, val_loss_list, names)

Model Channels: [16, 32, 48] Losses      Model Channels: [32, 48, 64] Losses      Model Channels: [64, 96, 128] Losses
Accuracy vs Epoch
```

可以看出，feature map nums对模型有一定的影响，我们选择64-96-128，即最大的。

3.探究 dense layer 对模型影响

模型定义

```
In [16]: class SimpleCNN(nn.Module):
    def __init__(self, dense_size=None):
        super(SimpleCNN, self).__init__()
        layers = []
        in_channels = 3

        # 固定的快速通道配置: [64, 96, 128]
        channels_config = [64, 96, 128]

        # 根据固定的快速通道配置设置卷积层
        for i in range(3):
            out_channels = channels_config[i]
            layers.append(nn.Conv2d(in_channels, out_channels, kernel_size=5, stride=1, padding=2))
            layers.append(nn.ReLU())
            layers.append(nn.MaxPool2d(kernel_size=2, stride=2))
            in_channels = out_channels

        # Flatten卷积后的输出
        layers.append(nn.Flatten())

        # 固定Dense层的配置: dense_size=[0-32-64-128-256到1024]
        if dense_size is None:
            dense_size = 32 # 默认配置为32

        layers.append(nn.Linear(in_channels * 16, dense_size)) # 16是经过池化后的每个特征图的大小
        layers.append(nn.ReLU())

        # 最后是10个分类
        layers.append(nn.Linear(dense_size, 10))

        self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)
```

实验过程

```
In [17]: # 实验过程
dense_sizes = [32, 64, 128, 256, 1024] # Dense层的尺寸变化配置
names = [f'Dense: {size}' for size in dense_sizes]
styles = ['-', '-', '-', '-', '-', '-', '-', '-', '-']

train_loss_list = []
train_acc_list = []
val_loss_list = []
val_acc_list = []

In [18]: # 进行实验
for dense_size in dense_sizes:
    print(f'Training with dense size {dense_size}...')
    model = SimpleCNN(dense_size=dense_size).to(device)
    train_loss, train_acc, val_loss, val_acc = train_model(model, train_loader, test_loader, epochs=20)
    train_loss_list.append(train_loss)
    train_acc_list.append(train_acc)
    val_loss_list.append(val_loss)
    val_acc_list.append(val_acc)

Training with dense size 32...
Epoch 1/10, Train Loss: 1.3736, Train Accuracy: 0.4979, Validation Loss: 1.0646, Validation Accuracy: 0.6220
Epoch 3/10, Train Loss: 0.7280, Train Accuracy: 0.7463, Validation Loss: 0.8746, Validation Accuracy: 0.7001
Epoch 5/10, Train Loss: 0.4925, Train Accuracy: 0.8269, Validation Loss: 0.7757, Validation Accuracy: 0.7309
Epoch 7/10, Train Loss: 0.3154, Train Accuracy: 0.8890, Validation Loss: 0.8637, Validation Accuracy: 0.7513
Epoch 9/10, Train Loss: 0.2039, Train Accuracy: 0.9284, Validation Loss: 1.0075, Validation Accuracy: 0.7406
Epoch 10/10, Train Loss: 0.1736, Train Accuracy: 0.9386, Validation Loss: 1.1163, Validation Accuracy: 0.7471
Training with dense size 128...
Epoch 1/10, Train Loss: 1.3708, Train Accuracy: 0.5017, Validation Loss: 1.0513, Validation Accuracy: 0.6236
Epoch 3/10, Train Loss: 0.7133, Train Accuracy: 0.7493, Validation Loss: 0.7721, Validation Accuracy: 0.7375
Epoch 5/10, Train Loss: 0.4666, Train Accuracy: 0.8354, Validation Loss: 0.8002, Validation Accuracy: 0.7356
Epoch 7/10, Train Loss: 0.2932, Train Accuracy: 0.8943, Validation Loss: 0.9608, Validation Accuracy: 0.7260
Epoch 9/10, Train Loss: 0.1913, Train Accuracy: 0.9317, Validation Loss: 1.0809, Validation Accuracy: 0.7432
Epoch 10/10, Train Loss: 0.1591, Train Accuracy: 0.9446, Validation Loss: 1.1719, Validation Accuracy: 0.7470
Training with dense size 256...
Epoch 1/10, Train Loss: 1.3918, Train Accuracy: 0.4919, Validation Loss: 1.0818, Validation Accuracy: 0.6114
Epoch 3/10, Train Loss: 0.7159, Train Accuracy: 0.7354, Validation Loss: 0.8365, Validation Accuracy: 0.7159
Epoch 5/10, Train Loss: 0.5091, Train Accuracy: 0.8198, Validation Loss: 0.7824, Validation Accuracy: 0.7400
Epoch 7/10, Train Loss: 0.3358, Train Accuracy: 0.8816, Validation Loss: 0.9143, Validation Accuracy: 0.7409
Epoch 9/10, Train Loss: 0.2179, Train Accuracy: 0.9225, Validation Loss: 1.0651, Validation Accuracy: 0.7484
Epoch 10/10, Train Loss: 0.1834, Train Accuracy: 0.9347, Validation Loss: 1.2453, Validation Accuracy: 0.7427
Training with dense size 1024...
Epoch 1/10, Train Loss: 1.3677, Train Accuracy: 0.5088, Validation Loss: 1.1876, Validation Accuracy: 0.6080
Epoch 3/10, Train Loss: 0.7417, Train Accuracy: 0.7437, Validation Loss: 0.8242, Validation Accuracy: 0.7126
Epoch 5/10, Train Loss: 0.4509, Train Accuracy: 0.8403, Validation Loss: 0.8316, Validation Accuracy: 0.7371
Epoch 7/10, Train Loss: 0.2719, Train Accuracy: 0.9164, Validation Loss: 1.0117, Validation Accuracy: 0.7371
Epoch 9/10, Train Loss: 0.1334, Train Accuracy: 0.9533, Validation Loss: 1.2565, Validation Accuracy: 0.7372
Epoch 10/10, Train Loss: 0.1347, Train Accuracy: 0.9544, Validation Loss: 1.3069, Validation Accuracy: 0.7342
```

实验结果

```
In [19]: # 绘制训练验证曲线
plot_val_acc(val_acc_list, names, styles=styles, title='Effect of Dense Layer Size on CIFAR-10')

Effect of Dense Layer Size on CIFAR-10
```

```
In [20]: # 绘制训练和验证损失曲线
plot_model_losses(train_loss_list, val_loss_list, names)

Model Dense: 32 Losses      Model Dense: 64 Losses      Model Dense: 128 Losses      Model Dense: 256 Losses      Model Dense: 1024 Losses
Accuracy vs Epoch
```

4.探究 How much dropout 对模型影响

模型定义

```
In [21]: class SimpleCNN(nn.Module):
    def __init__(self, dropout_rate=None):
        super(SimpleCNN, self).__init__()
        layers = []
        in_channels = 3

        # 固定的快速通道配置: [64, 96, 128]
        channels_config = [64, 96, 128]

        # 根据固定的快速通道配置设置卷积层
        for i in range(3):
            out_channels = channels_config[i]
            layers.append(nn.Conv2d(in_channels, out_channels, kernel_size=5, stride=1, padding=2))
            layers.append(nn.ReLU())
            layers.append(nn.MaxPool2d(kernel_size=2, stride=2))
            in_channels = out_channels

        # Flatten卷积后的输出
        layers.append(nn.Flatten())

        # 固定Dense层的配置 = 256
        layers.append(nn.Linear(in_channels * 16, 256)) # 16是经过池化后的每个特征图的大小
        layers.append(nn.ReLU())

        # 如果指定了 dropout 率，则添加 dropout 层
        if dropout_rate is not None:
            layers.append(nn.Dropout(dropout_rate))

        # 最后是10个分类
        layers.append(nn.Linear(256, 10))

        self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)
```

实验过程

```
In [22]: # 实验过程
dropout_rates = [0.1, 0.2, 0.3, 0.4, 0.5] # Dropout层的衰减配置
names = [f'Dropout: {rate}' for rate in dropout_rates]
styles = ['-', '-', '-', '-', '-', '-', '-', '-', '-']

train_loss_list = []
train_acc_list = []
val_loss_list = []
val_acc_list = []

In [23]: # 进行实验
for dropout_rate in dropout_rates:
    print(f'Training with dropout rate {dropout_rate}...')
    model = SimpleCNN(dropout_rate=dropout_rate).to(device)
    train_loss, train_acc, val_loss, val_acc = train_model(model, train_loader, test_loader, epochs=20)
    train_loss_list.append(train_loss)
    train_acc_list.append(train_acc)
    val_loss_list.append(val_loss)
    val_acc_list.append(val_acc)

Training with dropout rate 0.1...
Epoch 1/10, Train Loss: 1.3736, Train Accuracy: 0.4979, Validation Loss: 1.0646, Validation Accuracy: 0.6220
Epoch 3/10, Train Loss: 0.7280, Train Accuracy: 0.7463, Validation Loss: 0.8746, Validation Accuracy: 0.7001
Epoch 5/10, Train Loss: 0.4925, Train Accuracy: 0.8269, Validation Loss: 0.7757, Validation Accuracy: 0.7309
Epoch 7/10, Train Loss: 0.3154, Train Accuracy: 0.8890, Validation Loss: 0.8637, Validation Accuracy: 0.7513
Epoch 9/10, Train Loss: 0.2039, Train Accuracy: 0.9284, Validation Loss: 1.0075, Validation Accuracy: 0.7406
Epoch 10/10, Train Loss: 0.1736, Train Accuracy: 0.9386, Validation Loss: 1.1163, Validation Accuracy: 0.7471
Training with dropout rate 0.2...
Epoch 1/10, Train Loss: 1.3708, Train Accuracy: 0.5017, Validation Loss: 1.0513, Validation Accuracy: 0.6236
Epoch 3/10, Train Loss: 0.7133, Train Accuracy: 0.7493, Validation Loss: 0.7721, Validation Accuracy: 0.7375
Epoch 5/10, Train Loss: 0.4666, Train Accuracy: 0.8354, Validation Loss: 0.8002, Validation Accuracy: 0.7356
Epoch 7/10, Train Loss: 0.2932, Train Accuracy: 0.8943, Validation Loss: 0.9608, Validation Accuracy: 0.7260
Epoch 9/10, Train Loss: 0.1913, Train Accuracy: 0.9317, Validation Loss: 1.0809, Validation Accuracy: 0.7432
Epoch 10/10, Train Loss: 0.1591, Train Accuracy: 0.9446, Validation Loss: 1.1719, Validation Accuracy: 0.7470
Training with dropout rate 0.3...
Epoch 1/10, Train Loss: 1.3918, Train Accuracy: 0.4919, Validation Loss: 1.0818, Validation Accuracy: 0.6114
Epoch 3/10, Train Loss: 0.7159, Train Accuracy: 0.7354, Validation Loss: 0.8365, Validation Accuracy: 0.7159
Epoch 5/10, Train Loss: 0.5091, Train Accuracy: 0.8198, Validation Loss: 0.7824, Validation Accuracy: 0.7400
Epoch 7/10, Train Loss: 0.3358, Train Accuracy: 0.8816, Validation Loss: 0.9143, Validation Accuracy: 0.7409
Epoch 9/10, Train Loss: 0.2179, Train Accuracy: 0.9225, Validation Loss: 1.0651, Validation Accuracy: 0.7484
Epoch 10/10, Train Loss: 0.1834, Train Accuracy: 0.9347, Validation Loss: 1.2453, Validation Accuracy: 0.7427
Training with dropout rate 0.4...
Epoch 1/10, Train Loss: 1.3677, Train Accuracy: 0.5088, Validation Loss: 1.1876, Validation Accuracy: 0.6080
Epoch 3/10, Train Loss: 0.7417, Train Accuracy: 0.7437, Validation Loss: 0.8242, Validation Accuracy: 0.7126
Epoch 5/10, Train Loss: 0.4509, Train Accuracy: 0.8403, Validation Loss: 0.8316, Validation Accuracy: 0.7371
Epoch 7/10, Train Loss: 0.2719, Train Accuracy: 0.9164, Validation Loss: 1.0117, Validation Accuracy: 0.7371
Epoch 9/10, Train Loss: 0.1334, Train Accuracy: 0.9533, Validation Loss: 1.
```


Training with dropout rate 0.1...

Epoch 1/10, Train Loss: 1.3614, Train Accuracy: 0.5869, Validation Loss: 1.8579, Validation Accuracy: 0.6291

Epoch 3/10, Train Loss: 0.7117, Train Accuracy: 0.7520, Validation Loss: 0.8121, Validation Accuracy: 0.7270

Epoch 5/10, Train Loss: 0.4681, Train Accuracy: 0.8345, Validation Loss: 0.6251, Validation Accuracy: 0.7473

Epoch 7/10, Train Loss: 0.3868, Train Accuracy: 0.8986, Validation Loss: 0.7620, Validation Accuracy: 0.7595

Epoch 9/10, Train Loss: 0.2104, Train Accuracy: 0.9247, Validation Loss: 1.0436, Validation Accuracy: 0.7537

Epoch 10/10, Train Loss: 0.1781, Train Accuracy: 0.9373, Validation Loss: 1.1021, Validation Accuracy: 0.7547

Training with dropout rate 0.2...

Epoch 1/10, Train Loss: 1.4883, Train Accuracy: 0.4889, Validation Loss: 1.1835, Validation Accuracy: 0.6073

Epoch 3/10, Train Loss: 0.7703, Train Accuracy: 0.7331, Validation Loss: 0.7829, Validation Accuracy: 0.7308

Epoch 5/10, Train Loss: 0.5409, Train Accuracy: 0.8111, Validation Loss: 0.7555, Validation Accuracy: 0.7459

Epoch 7/10, Train Loss: 0.4632, Train Accuracy: 0.8632, Validation Loss: 0.7613, Validation Accuracy: 0.7492

Epoch 9/10, Train Loss: 0.2757, Train Accuracy: 0.9813, Validation Loss: 0.8712, Validation Accuracy: 0.7577

Epoch 10/10, Train Loss: 0.2420, Train Accuracy: 0.9132, Validation Loss: 0.9894, Validation Accuracy: 0.7552

Training with dropout rate 0.3...

Epoch 1/10, Train Loss: 1.4474, Train Accuracy: 0.4699, Validation Loss: 1.1620, Validation Accuracy: 0.5854

Epoch 3/10, Train Loss: 0.8336, Train Accuracy: 0.7084, Validation Loss: 0.8603, Validation Accuracy: 0.7033

Epoch 5/10, Train Loss: 0.6890, Train Accuracy: 0.7860, Validation Loss: 0.7576, Validation Accuracy: 0.7420

Epoch 7/10, Train Loss: 0.4787, Train Accuracy: 0.8365, Validation Loss: 0.7613, Validation Accuracy: 0.7579

Epoch 9/10, Train Loss: 0.3581, Train Accuracy: 0.8718, Validation Loss: 0.8225, Validation Accuracy: 0.7476

Epoch 10/10, Train Loss: 0.3166, Train Accuracy: 0.8857, Validation Loss: 0.9037, Validation Accuracy: 0.7600

Training with dropout rate 0.4...

Epoch 1/10, Train Loss: 1.4778, Train Accuracy: 0.4680, Validation Loss: 1.1343, Validation Accuracy: 0.5869

Epoch 3/10, Train Loss: 0.8409, Train Accuracy: 0.7865, Validation Loss: 0.8045, Validation Accuracy: 0.7238

Epoch 5/10, Train Loss: 0.6233, Train Accuracy: 0.7838, Validation Loss: 0.7622, Validation Accuracy: 0.7469

Epoch 7/10, Train Loss: 0.4783, Train Accuracy: 0.8327, Validation Loss: 0.8340, Validation Accuracy: 0.7424

Epoch 9/10, Train Loss: 0.3659, Train Accuracy: 0.8789, Validation Loss: 0.8998, Validation Accuracy: 0.7483

Epoch 10/10, Train Loss: 0.3227, Train Accuracy: 0.8873, Validation Loss: 0.8650, Validation Accuracy: 0.7567

Training with dropout rate 0.5...

Epoch 1/10, Train Loss: 1.4708, Train Accuracy: 0.4587, Validation Loss: 1.1431, Validation Accuracy: 0.5906

Epoch 3/10, Train Loss: 0.8725, Train Accuracy: 0.6986, Validation Loss: 0.8259, Validation Accuracy: 0.7136

Epoch 5/10, Train Loss: 0.6634, Train Accuracy: 0.7690, Validation Loss: 0.7487, Validation Accuracy: 0.7443

Epoch 7/10, Train Loss: 0.5328, Train Accuracy: 0.8139, Validation Loss: 0.7603, Validation Accuracy: 0.7451

Epoch 9/10, Train Loss: 0.4345, Train Accuracy: 0.8461, Validation Loss: 0.8624, Validation Accuracy: 0.7572

Epoch 10/10, Train Loss: 0.3855, Train Accuracy: 0.8626, Validation Loss: 0.8313, Validation Accuracy: 0.7511

实验结果



5. 加入 Augmentation 对模型影响

数据增强

In [25]: # 强力的数据增强组合，只应用于训练集

```
train_transform = transforms.Compose([
    transforms.RandomHorizontalFlip(), # 随机水平翻转，保持不变
    transforms.RandomRotation(5), # 随机旋转角度减小，从15度减小到5度
    transforms.RandomResizedCrop(32, scale=(0.9, 1.0)), # 随机裁剪，缩放范围减小（不小于90%）
    transforms.ColorJitter(brightness=0.1, contrast=0.1, saturation=0.1, hue=0.1), # 随机调整颜色幅度
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)) # 归一化
])

# 仅进行标准化的转换，应用于测试集
test_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)) # 归一化
])

# 加载CIFAR-10数据集
train_dataset_aug = datasets.CIFAR10(root="./data", train=True, download=True, transform=train_transform)
train_dataset = datasets.CIFAR10(root="./data", train=False, download=True, transform=test_transform)

train_loader_aug = DataLoader(train_dataset_aug, batch_size=64, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)
```

Files already downloaded and verified
Files already downloaded and verified

模型定义

In [26]:

```
for images, labels in train_loader_aug:
    # 选择第一张图像
    image = labels[2]

    # 将图像从Tensor格式转换为NumPy格式
    image = image.numpy().transpose((1, 2, 0)) # 由(3, 32, 32) -> (32, 32, 3)
    image = image * 0.5 + 0.5 # 反标准化，将像素值从[-1, 1]映射回[0, 1]

    # 显示图像
    plt.imshow(image)
    plt.title(f"Label: {labels[0].item()}")
    plt.axis('off') # 不显示坐标轴
    plt.show()
    break # 只显示一个批次中的第一张图像
```



In [27]:

```
class SimpleCNN(nn.Module):
    def __init__(self):
        super(SimpleCNN, self).__init__()
        layers = []
        in_channels = 3

        # 固定的通道数配置: [64, 96, 128]
        channels_config = [64, 96, 128]

        # 根据固定的通道数配置设置卷积层
        for i in range(3):
            out_channels = channels_config[i]
            layers.append(nn.Conv2d(in_channels, out_channels, kernel_size=5, stride=1, padding=2))
            layers.append(nn.MaxPool2d(kernel_size=2, stride=2))
            in_channels = out_channels

        # Flatten卷积后的输出
        layers.append(nn.Flatten())

        # 固定Dense层的配置: 256
        layers.append(nn.Linear(in_channels * 16, 256)) # 16是经过池化后的每个特征图的大小
        layers.append(nn.ReLU())

        layers.append(nn.Dropout(0.5))

        # 最后是10个分类
        layers.append(nn.Linear(256, 10))

        self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)
```

实验过程

In [28]:

```
model = SimpleCNN().to(device)
# 运行训练
train_loss, train_acc, val_loss, val_acc = train_model(model, train_loader_aug, test_loader, epochs=30)
```

Epoch 1/30, Train Loss: 1.5701, Train Accuracy: 0.4244, Validation Loss: 1.2412, Validation Accuracy: 0.6524

Epoch 3/30, Train Loss: 1.0553, Train Accuracy: 0.6336, Validation Loss: 0.9801, Validation Accuracy: 0.6813

Epoch 5/30, Train Loss: 0.8995, Train Accuracy: 0.6918, Validation Loss: 0.8107, Validation Accuracy: 0.7278

Epoch 7/30, Train Loss: 0.8181, Train Accuracy: 0.7182, Validation Loss: 0.7947, Validation Accuracy: 0.7265

Epoch 9/30, Train Loss: 0.7621, Train Accuracy: 0.7373, Validation Loss: 0.7275, Validation Accuracy: 0.7526

Epoch 11/30, Train Loss: 0.7229, Train Accuracy: 0.7518, Validation Loss: 0.7201, Validation Accuracy: 0.7581

Epoch 13/30, Train Loss: 0.6912, Train Accuracy: 0.7632, Validation Loss: 0.6923, Validation Accuracy: 0.7693

Epoch 15/30, Train Loss: 0.6690, Train Accuracy: 0.7724, Validation Loss: 0.6676, Validation Accuracy: 0.7781

Epoch 17/30, Train Loss: 0.6485, Train Accuracy: 0.7784, Validation Loss: 0.6705, Validation Accuracy: 0.7798

Epoch 19/30, Train Loss: 0.6246, Train Accuracy: 0.7850, Validation Loss: 0.6646, Validation Accuracy: 0.7884

Epoch 21/30, Train Loss: 0.6048, Train Accuracy: 0.7908, Validation Loss: 0.6932, Validation Accuracy: 0.7851

Epoch 23/30, Train Loss: 0.5996, Train Accuracy: 0.7942, Validation Loss: 0.6810, Validation Accuracy: 0.7831

Epoch 25/30, Train Loss: 0.5825, Train Accuracy: 0.8027, Validation Loss: 0.6530, Validation Accuracy: 0.7937

Epoch 27/30, Train Loss: 0.5758, Train Accuracy: 0.8021, Validation Loss: 0.6441, Validation Accuracy: 0.7830

Epoch 29/30, Train Loss: 0.5552, Train Accuracy: 0.8091, Validation Loss: 0.6364, Validation Accuracy: 0.7926

Epoch 30/30, Train Loss: 0.5608, Train Accuracy: 0.8079, Validation Loss: 0.6613, Validation Accuracy: 0.7885



实验结果

我们可以看出，加上数据增强之后模型的泛化能力极大提高，上升了5%左右

实验总结

我们可以看出，数据增强对模型的泛化能力提升最强，约为5%左右。此外，特征图通道数量的深度对模型的提升也非常大。即使我们用了非常小的模型（仅仅三层），用上比较好的调参技巧，也可以在 cifar10 上把模型训练到80%的准确率。通过这次实验，我们能够很好的理解各种参数对模型的作用。